

Project 1

Distributed Data Infrastructure

Walter Genchi 014961054

November 27, 2018

1 Problem 1

The idea behind this exercise is to divide the data in bags (`n_of_bags` is the a-priori parameter to be chosen) and determine the bag where the median is located. The task will be thus to sort only that bag.

There is a trade-off in the choice of `n_of_bags`. From one hand, for large `n_of_bags`, the sorting task on the median bag (see step 8) will be simplified, but the Spark operations (see step 5 and 6) will be affected by an increased number of keys. On the other hand, for small `n_of_bags`, the sorting task will be slower and Spark operations faster.

Moreover, the bag-assigning strategy (see step 3) is effective for uniformly distributed data (i.e. all bags have a similar number of elements), but lacks generality for other data distributions (e.g. highly skewed distributions).

Operations can be summarized as:

1. Reading the file `data-1.txt` and converting the numbers to float (using the `map` function)
2. Make data persistent, in order to make further actions on data (i.e. `max`, `min` and `count`) more efficient in terms of time and memory consumption.
3. Every number n is mapped to his bag (using the `map` function) by 3 steps:
 - Normalization, i.e. $x = (n - \min) / (\max - \min)$
 - Assigned bag = $\text{int}(x * \text{n_of_bags})$
 - FINAL RESULT: $n \rightarrow (\text{Assigned Bag}, n)$
4. Computing the median position (by definition of median the value in the middle)
5. Creating the bags, i.e. get something like

`[(0, [numbers in 0 - 0.0999]), ..., (10, numbers in 0.9 - 1)]`

by using `groupByKey`. Note that `reduceByKey` is usually preferred for grouping + aggregating tasks (e.g. word count) due to better performances on large dataset. However, since we require the grouped bags for future operations, we perform first the grouping (`groupByKey`) and then the size (`mapValues(len)`) of the bags.

6. Computing the bag size, i.e. get something like

`[(0, [how many numbers in 0 - 0.0999]), ..., (10, how many numbers in 0.9 - 1)]`

by using `mapValues(len)`.

7. Determining the bag which contains the median, by computing cumulative sums of the size of each bag, starting from the first bag.
8. Sorting the median bag and writing the result.

2 Problem 2

The points I took into account while considering how to solve this problem were the following:

- A has much more rows (10^6) than columns (10^3).
- The product between three matrices has the associative property, i.e. $(AB)C = A(BC)$. Thus, I have preferred to perform first $A^T A$, since the result is a $10^3 \times 10^3$ matrix, which is easier to store in memory than the AA^T matrix, which has size $10^6 \times 10^3$.
- While using `map` function in Spark, I noticed that it is easy to iterate over RDD rows (by using a `lambda` function), but it is hard to iterate over RDD columns (unless you transpose the matrix, which is probably not feasible because of its dimension).

Compute $B = A^T A$

In order to use `map` (transformation) and `reduce` (action) in Spark, I exploited the property of outer product in matrices.

Given matrix

$$A = \begin{pmatrix} a & b \\ c & d \\ e & f \end{pmatrix}$$

then

$$B = A^T A = \begin{pmatrix} a & c & e \\ b & d & f \end{pmatrix} \begin{pmatrix} a & b \\ c & d \\ e & f \end{pmatrix} = \begin{pmatrix} a^2 + c^2 + e^2 & ab + cd + ef \\ ba + dc + fe & b^2 + d^2 + f^2 \end{pmatrix}.$$

This product can be computed as the sum of the outer products of the rows of A , i.e.

$$\begin{aligned} \begin{pmatrix} a & b \end{pmatrix} \otimes \begin{pmatrix} a & b \end{pmatrix} + \dots + \begin{pmatrix} e & f \end{pmatrix} \otimes \begin{pmatrix} e & f \end{pmatrix} = \\ \begin{pmatrix} a^2 & ab \\ ba & b^2 \end{pmatrix} + \dots + \begin{pmatrix} e^2 & ef \\ fe & f^2 \end{pmatrix} = \\ \begin{pmatrix} a^2 + c^2 + e^2 & ab + cd + ef \\ ba + dc + fe & b^2 + d^2 + f^2 \end{pmatrix} = A^T A \end{aligned}$$

In Spark this means that the operation (`np.outer`) is always computed on the same row and thus very efficient with `map`.

After that, the intermediate result is one RDD, where each element k is a `np.array` of size $10^3 \times 10^3$, obtained as `np.outer(k,k)`. This is the input of the reduce function, which adds the 10^6 matrices (each with dimension $10^3 \times 10^3$) through the commutative operator `add`.

This solution works well with the sample dataset (10^3 rows: ~ 1 second for `map` and ~ 3 seconds for `reduce`) and with a bigger sample (10^5 rows: ~ 2 seconds for `map` and ~ 56 seconds for `reduce`), which I generated from the original dataset.

However, when testing on the whole matrix (10^6 rows), I got an error on the reduce operation. That is the reason why I tried to implement a different code, which first assigns an index to each element of the RDD, then maps each index to one key (the set of keys is "smaller" than the set of indices) and finally uses `reduceByKey` to add the matrices with the same key. This simplifies the computations in the reduce part and acts as a "combiner" in MapReduce. Finally, the results are again summed up using `reduce(add)`, this time with less matrices. This second solution did not provide satisfactory results on the whole matrix (10^6 rows), due to errors in the `reduce(add)` function.

Compute $W = AB$

This product was implemented (using `map` in Spark) by computing the numpy dot product between each row of A and the W matrix, which results as a 1×1000 vector. The dot product between the first row of A and matrix B constitutes the first row of the final matrix W . The results are finally sorted by key, since RDD does not guarantee the order of the computations in the output.